

Attributed Graph Clustering


Presented by Zhao Zhou

Outline

- Motivation
- Related work
- Attributed Graph
- Distance-based Clustering Algorithm
- Model-based Clustering Algorithm
- Conclusion

Graph Clustering

- For a given set of objects, we would like to divide it into groups of similar objects.
- The similarity between objects is defined by objective functions



Real Applications

- Community Detection
 - Nodes represent members
 - Friendship relationship represent the corresponding links
- Telecommunication Networks
 - Individual phone numbers represent nodes
 - Phone calls represented as edges
- Email Analysis
 - Individuals represented as nodes
 - Emails sent represented as edges

Related work

- A number of algorithms for graph clustering proposed before
 1. The Minimum Cut Algorithm [3]
 2. Multi-way Graph Partition [4]
 3. k-medoid and k-means algorithm [5]
 4. Spectral Clustering method [6]
 5.

Attributed Graph

- Existing graph clustering methods
 - mainly focus on the topological structure for clustering
 - Ignore the vertex properties
- Proliferation of rich information for real objects
 - Vertices associated with a number of attributes
 - Its attributes describe the characteristics and properties
 - give rise to a new type of graphs

Applications

- Attributed graph clustering is useful in many domains
 - Service-oriented social network
 - Communication network

Service-oriented social network

- Each user can be characterized by his profile
 - Interests, gender, education, etc
- Clustering users by considering both relationships and profile is more useful
 - Recommendation system
 - User-targeted online advertising

Communication network

- Each subscriber is associated with
 - Demographic information, service usage, etc
- Attributed Graph clustering on users is useful in
 - Design effective group-oriented marketing strategies

Problem Definition

- Given an attributed graph G and the number of clusters K , the clustering is to partition the vertex set V of G into K disjoint subsets V_1, \dots, V_k
 - Vertices within clusters are densely connected
 - Vertices in different clusters are sparsely connected
 - Vertices within clusters have low diversity in their attribute values
 - Vertices in different clusters have diverse attribute values


Attributed Graph Clustering Algorithm

- The algorithms for attributed graph clustering can be categorized into two types
 - Distance-based Approaches
 - Model-based Approaches

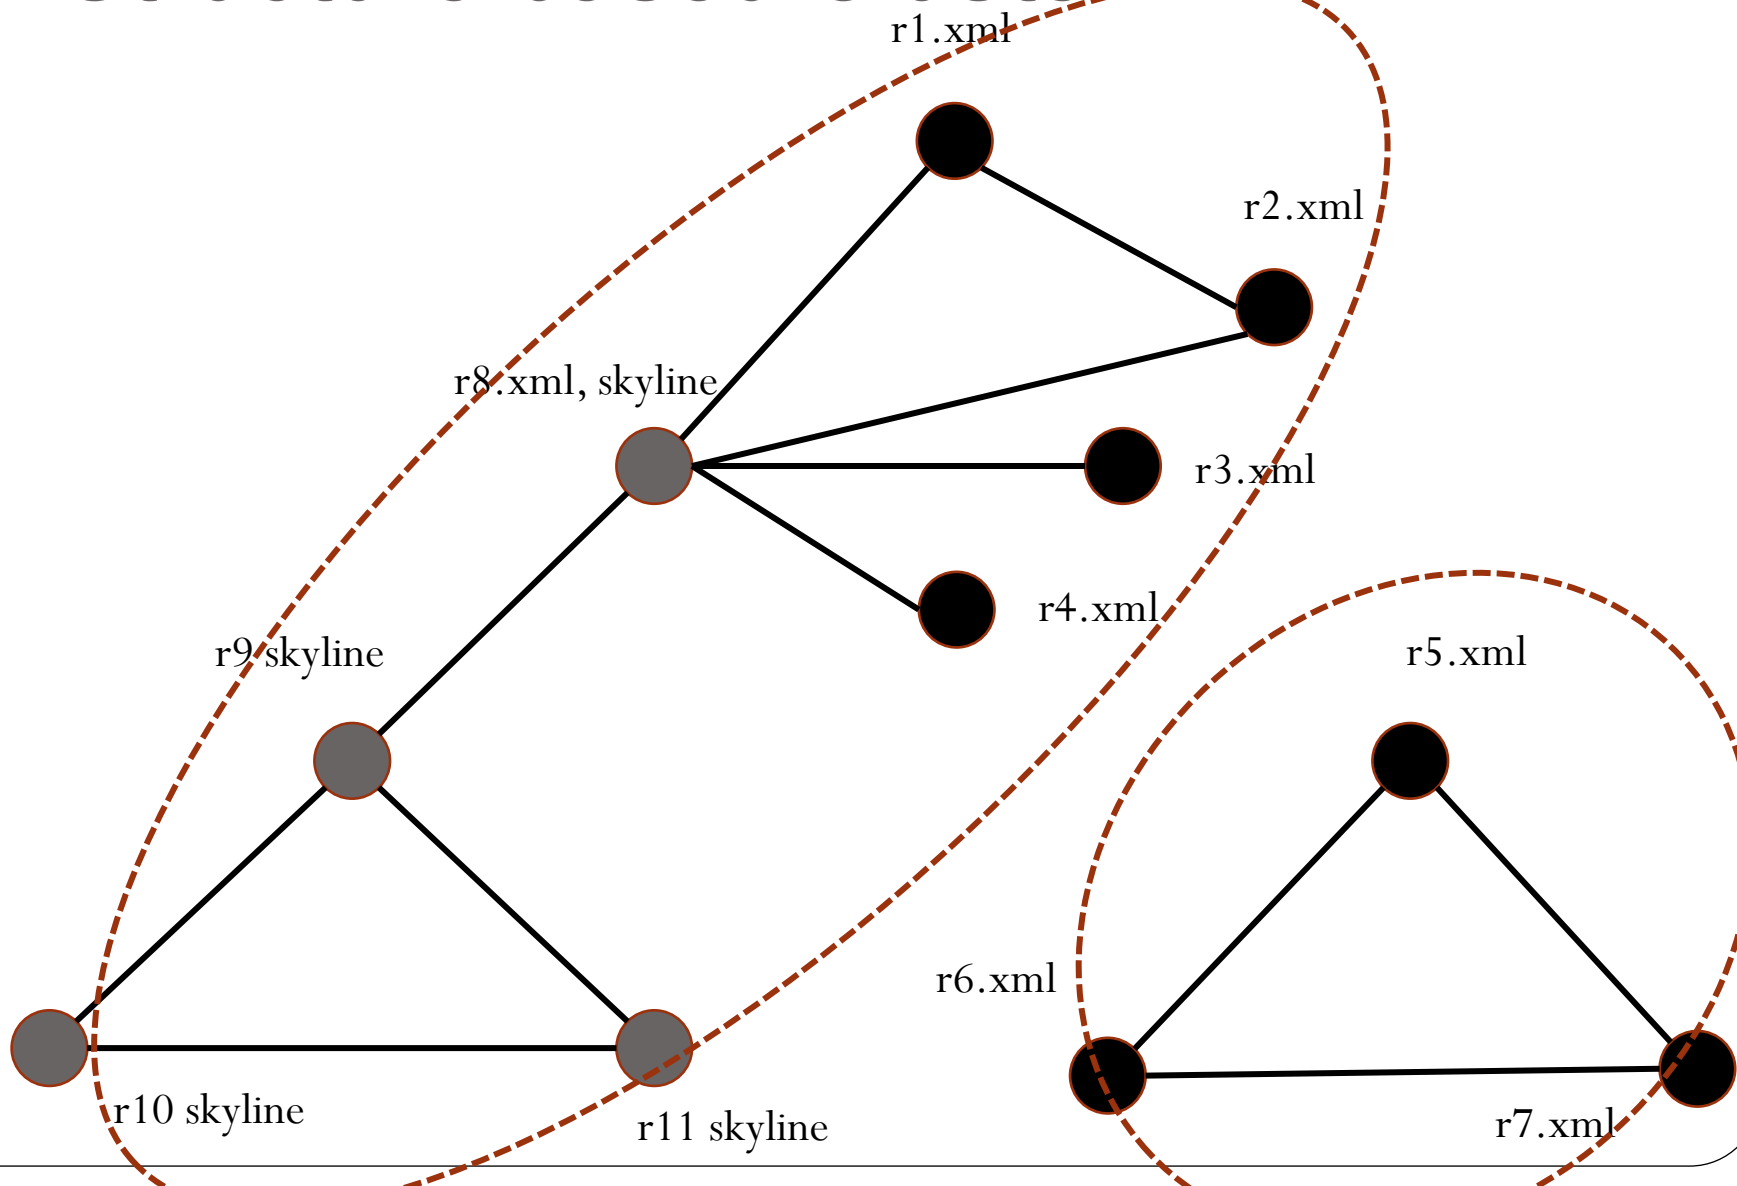
Distance-based Approaches(vldb09)

- The main idea is to design a distance measure for vertex pairs that combines both structural and attribute information
- Based on this measure, standard clustering algorithms like k-medoids are then applied to cluster the vertices
- The state-of-the-art approaches are the SA-cluster [1] and its extended version [7,8]


Co-author Graph



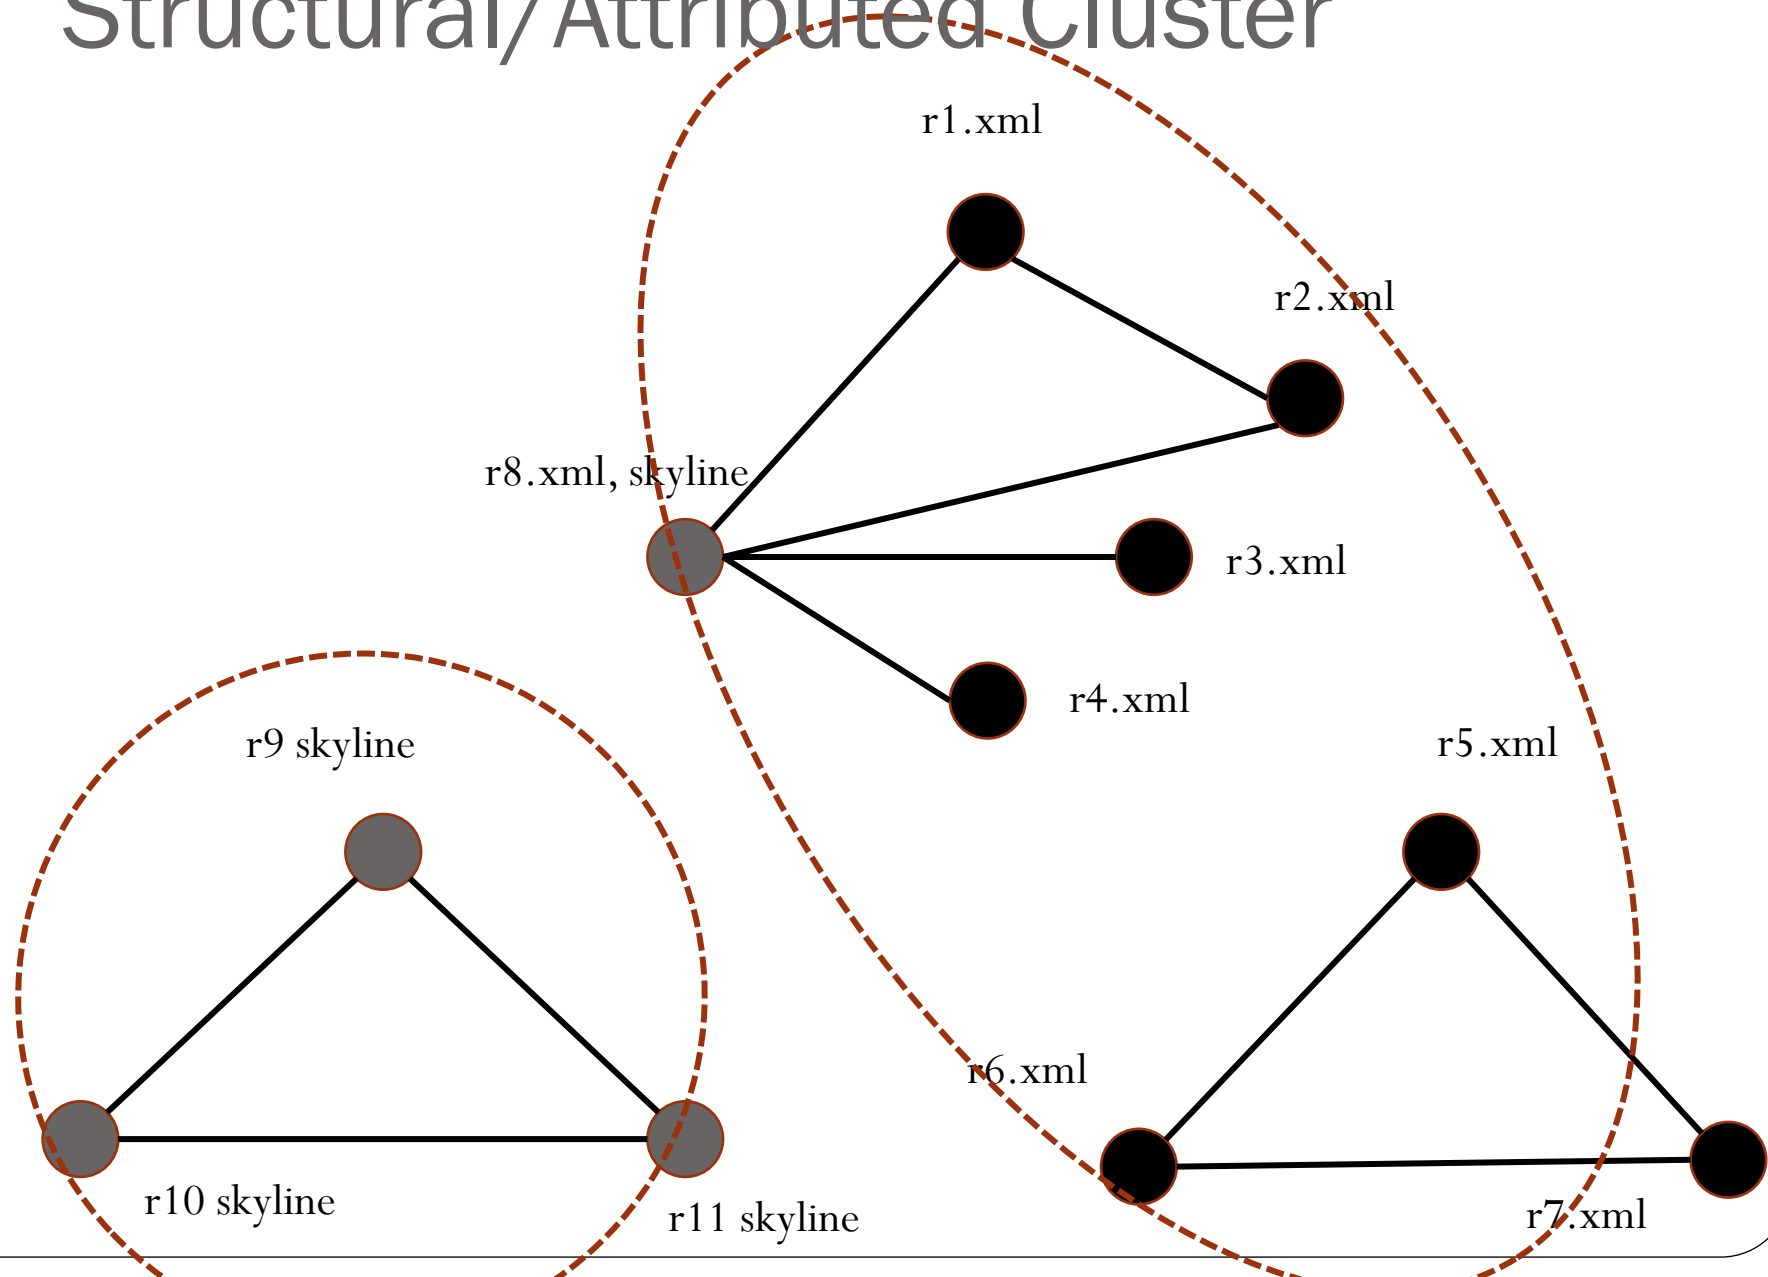
Structure-based Cluster



Attribute-based Cluster



Structural/Attributed Cluster



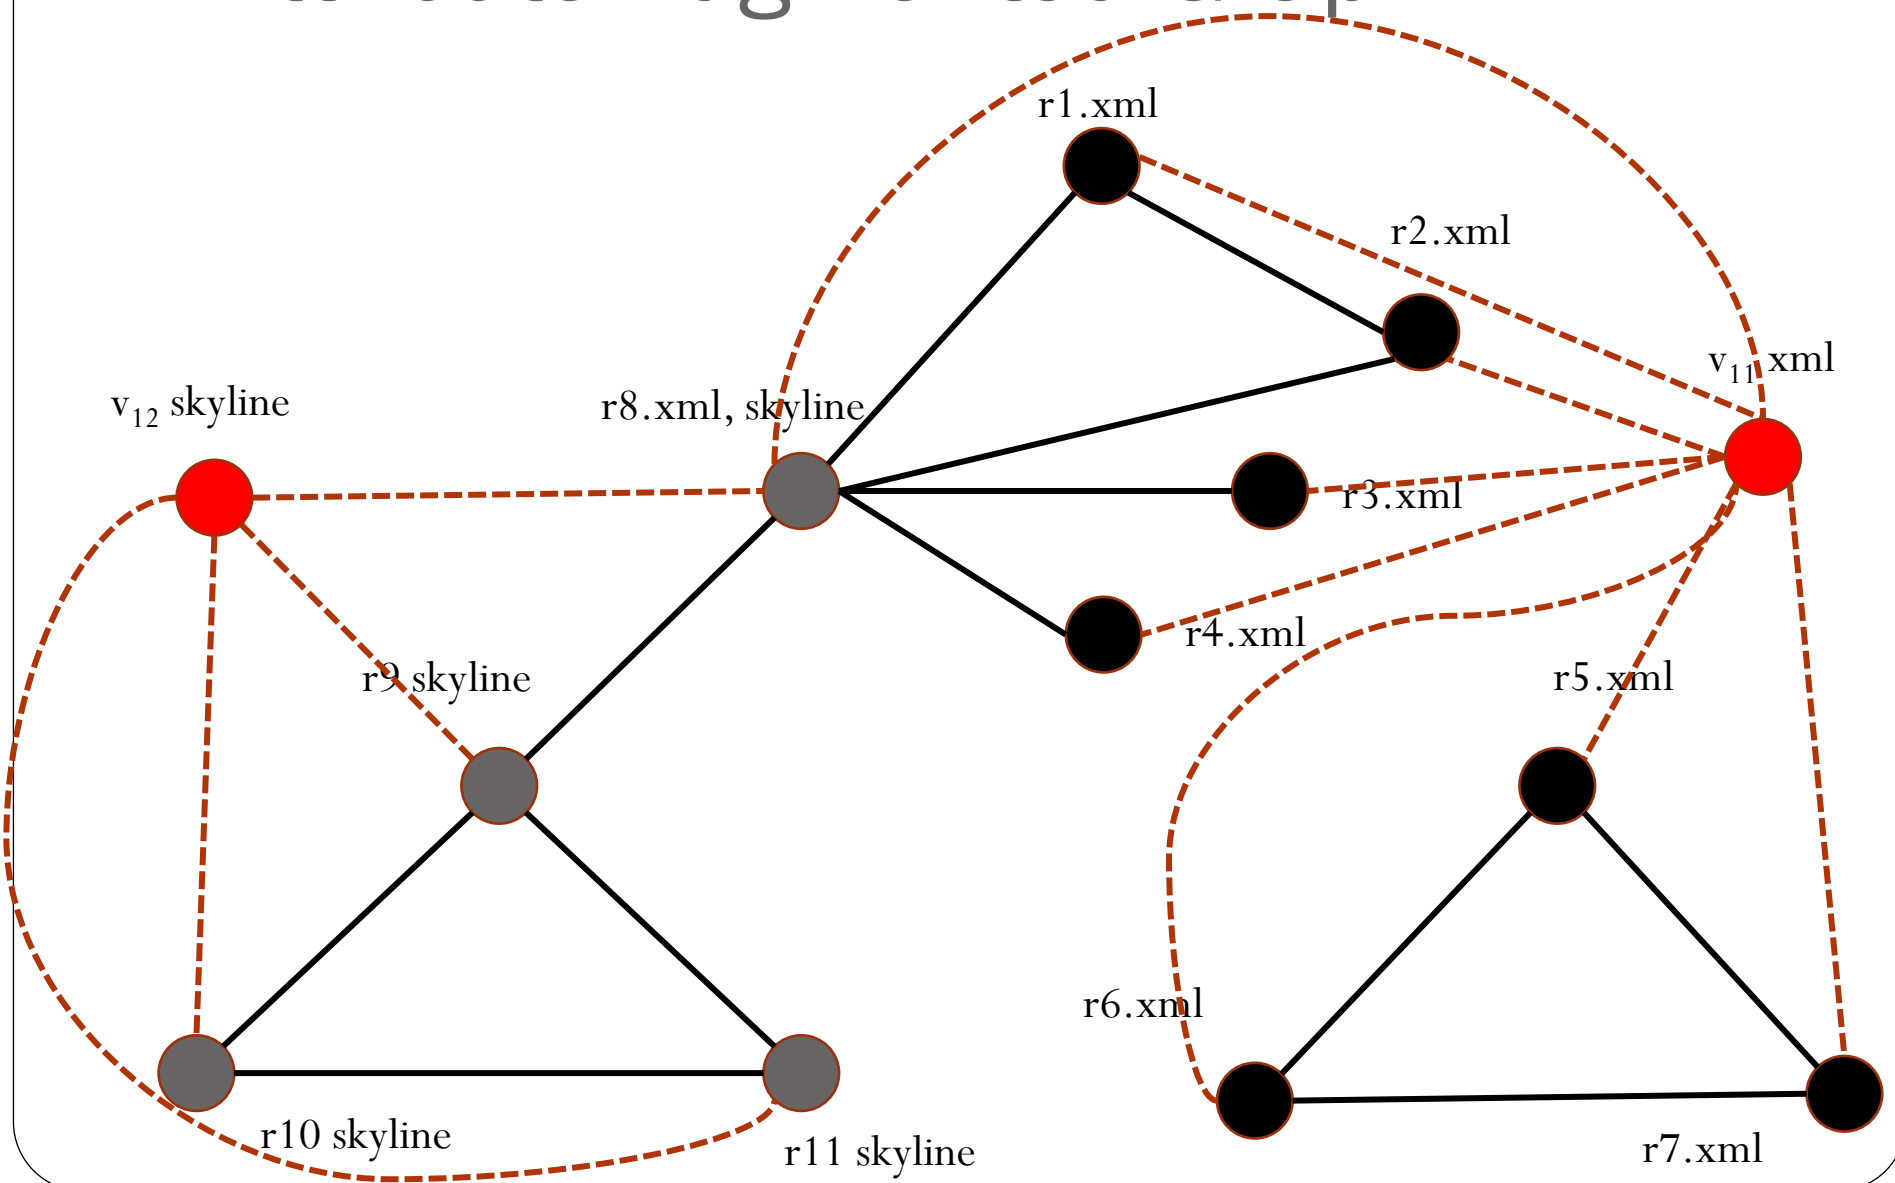
Basic solution

- Design a distance function between v_i and v_j
 - $d(v_i, v_j) = \alpha \cdot d_S(v_i, v_j) + \beta \cdot d_A(v_i, v_j)$
- Hard to set the parameters α and β

Attribute augmented Graph

- An attribute augmented graph is $G_a = (V \cup V_a, E \cup E_a)$
 - V_a is the set of attribute vertex, v_{ij} represents attribute i takes the j th value
 - E_a is attribute edge, (v_i, v_{ij}) means vertex v_i takes a value of a_{ij} on attribute i

Attribute Augmented Graph



Transition probability

weight of structure edge

- Transition probability between structure vertex

- $P(v_i, v_j) = \frac{w_0}{|N(v_i)| * w_0 + w_1 + \dots + w_m}$

weight of attribute edge

- Transition probability through structure vertex to attribute vertex

- $P(v_i, v_j) = \frac{w_j}{|N(v_i)| * w_0 + w_1 + \dots + w_m}$

- Transition probability through attribute vertex to structure vertex

- $P(v_{ik}, v_j) = \frac{1}{|N(v_{ik})|}$

Transition probability matrix

- Presentation of transition probability matrix
 - $P_A = \begin{bmatrix} P_V & A \\ B & o \end{bmatrix}$
 - P_V transition probability between vertices
 - A transition probability between vertices and attribute vertices
 - B transition probability between attribute vertices and vertices
- P_A can be incrementally computed by
 - $P_A^l = \begin{bmatrix} T_l & T_{l-1}A \\ BT_{l-1} & BT_{l-2}A \end{bmatrix}$
 - $T_l = P_V T_{l-1} + ABT_{l-2}$

Random Walk Distance

- Random walk distance
 - $R_A^l = \sum_{\tau=1}^l c(1-c)^\tau P_A^\tau$
 - τ length of path
 - c decay factor
 - P_A probability transition matrix

Initialization of Clustering Method

- Density Function

- $f(v_i) = \sum_{v_j \in V} (1 - e^{-\frac{d(v_i, v_j)^2}{2\sigma^2}})$
- σ is a user-specified parameter

- Select the densest k vertices on $f(v_i)$ as the initial centroids $\{c_1^0, \dots, c_k^0\}$

Clustering Process

- Assign vertex to its closest centroid c^*
 - $c^* = \operatorname{argmax}_{c_j^t} d(v_i, c_j^t)$
- Compute the “average point” of a cluster V_i
 - $R_A^l(\operatorname{avg}(v_i), v_j) = \frac{1}{|V_i|} \sum_{v_k \in V_i} R_A^l(v_k, v_j)$
- Find the new centroid c_i^{t+1} in cluster V_i
 - $c_i^{t+1} = \operatorname{argmin}_{v_j \in V_i} ||R_A^l(v_j) - R_A^l(\operatorname{avg}(v_i))||$

Weight Self-Adjustment

- If a large portion of vertices within clusters shares the same value of a certain attribute a_i , it means that a_i has a good clustering accuracy.
- If vertices within clusters have a very random distribution on values of a certain attribute a_i , then a_i is not a good clustering attribute.

Weight Self-Adjustment

- $\text{Vote}_i(v_p, v_q) = 1$ if v_p and v_q share the same value on a_i

-

$$\omega_i^{t+1} = \frac{1}{2}(\omega_i^t + \Delta\omega_i^t) = \frac{1}{2}\left(\omega_i^t + \frac{m \sum_{j=1}^k \sum_{v \in V_j} \text{vote}_i(c_j, v)}{\sum_{p=1}^m \sum_{j=1}^k \sum_{v \in V_j} \text{vote}_p(c_j, v)}\right)$$


Clustering Procedure

1. Assign vertices to their closest centroids
2. Update cluster centroids
3. Adjust attribute edge weights $\{w_1, \dots, w_m\}$
4. Re-calculate the random walk distance matrix R_A


Experimental Result

- SA-Cluster: considers both structural and attribute similarities
- S-Cluster: only considers structural distance
- W-Cluster: consider structural and attribute similarities equally important
- K-SNAP: partition vertices with the same attribute values into a cluster

Experimental Result



(a) density



(b) entropy

Model-based Approach(SIGMOD12)

- Avoids the artificial design of a distance measure
- Based on probabilistic model
 - Fuses both structural and attribute information
 - Vertices from the same cluster behave similar, while vertices from different clusters can behave differently.
- Clustering with the proposed model can be transformed into a probabilistic inference problem
 - Model enforces the intra-cluster similarity by asserting that the attribute values and edge connections of a vertex should depend on its cluster label

Problem Statement

- An attributed graph G is a 4-tuple (V, E, Λ, F)
 - $V = \{v_1, v_2, \dots, v_N\}$: a set of N vertices
 - $E = \{(v_i, v_j)\}$: a set of edges
 - $\Lambda = \{a^1, a^2, \dots, a^T\}$: a set of T categorical attributes
 - $F = \{f_1, f_2, \dots, f_T\}$: a set of functions map each element in attribute vector to attribute value i.e. $f_t(v_i)$

Bayesian Model for Clustering

- Given a set of vertices V , a set of attributes Λ , and the number of clusters K , the model defines a joint probability distribution over all possible partitions over V .
- The goal is to find the most probable clustering with the maximum probability.

Notations

- X : adjacency matrix ($N \times N$)
 - X_{ij} is a binary random variable
 - Indicates the existence of edge
- Y : attribute matrix ($N \times T$)
 - Y_i^t is a categorical random variable
 - Denotes the value of attribute a^t associated with vertex v_i
- Z : clustering of vertices ($N \times 1$)
 - Z_i is a categorical random variable
 - Denotes the label of vertex v_i


Generating Z

- Cluster label Z_i of each vertex v_i is sampled from multinomial distribution independently
 - $P(Z_i = k \mid \alpha) = \alpha_k, k = 1, 2, \dots, K$
 - α_k denotes the proportion of the vertices belonging to cluster k
 - $\sum_{k=1}^K \alpha_k = 1, \alpha_k \in [0, 1]$




Generating Y

- Given cluster label Z_i for vertex v_i , the attribute values of v_i is by
 - $P(Y_i^t = m \mid \theta_{Zi}^t) = \theta_{Zim}^t, m = 1, 2, \dots, M^t$
 - M^t is the size of the domain $\text{dom}(a^t)$
 - θ_{Zim}^t denotes the proportion of vertices in cluster Z_i that take the m -th value in $\text{dom}(a^t)$
 - $\sum_{m=1}^{M^t} \theta_{Zim}^t = 1, \theta_{Zim}^t \in [0, 1]$

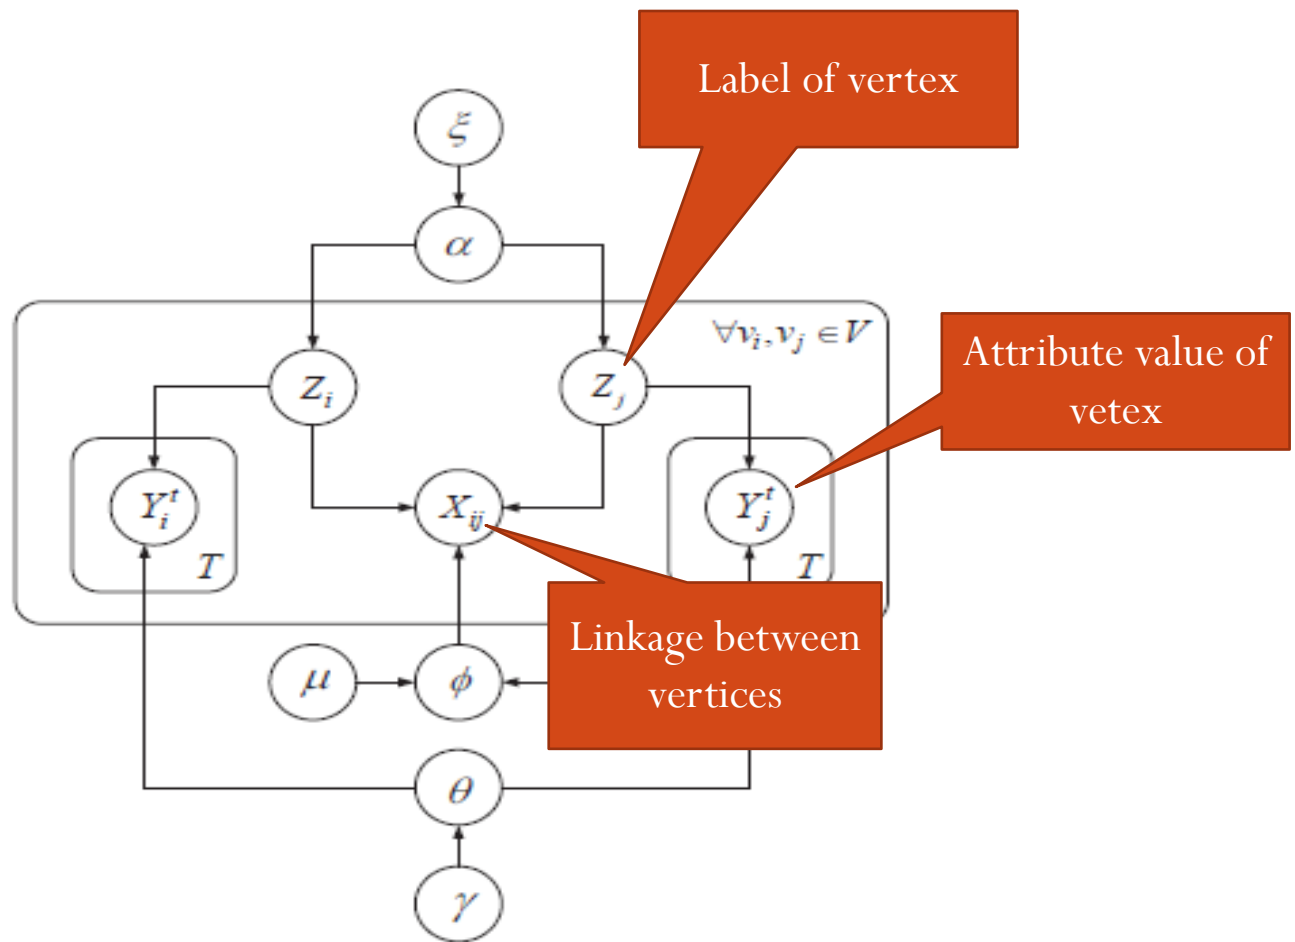


Generating X

- Given the cluster labels Z_i and Z_j for vertex v_i and v_j , the indicator X_{ij} denotes the existence of an edge between them
 - X_{ij} is a binary variable taking value 0 or 1
 - $P(X_{ij} | \phi_{Z_i Z_j}) = (1 - \phi_{Z_i Z_j})^{1 - X_{ij}} (\phi_{Z_i Z_j})^{X_{ij}}$
 - $\phi_{Z_i Z_j} \in [0, 1]$



Graphical representation



Model Definition

- Given the hyper-parameters $\epsilon, \gamma, \mu, \nu$, the joint distribution over $\alpha, \theta, \phi, X, Y, Z$ can be decomposed as

$$\begin{aligned} & p(\alpha, \theta, \phi, \mathbf{X}, \mathbf{Y}, \mathbf{Z} | \xi, \gamma, \mu, \nu) \\ = & p(\alpha | \xi) p(\theta | \gamma) p(\phi | \mu, \nu) p(\mathbf{Z} | \alpha) p(\mathbf{X} | \mathbf{Z}, \phi) p(\mathbf{Y} | \mathbf{Z}, \theta), \end{aligned}$$

Model-based Clustering of Attributed Graph

- The problem of clustering a given attributed graph (X, Y) can be transformed into a standard probabilistic inference problem
- Find the clustering Z conditioning on X, Y such that
 - $Z^* = \operatorname{argmax}_z p(Z | X, Y)$

Challenges

- The first difficulty is the maximization over the N variables $\mathbf{Z} = \{Z_1, Z_2, \dots, Z_N\}$. For large N , the global maximization is computationally prohibitive

- The second difficulty is in the calculation of the posterior distribution

$$p(\mathbf{Z}|\mathbf{X}, \mathbf{Y}) = \iiint p(\alpha, \theta, \phi, \mathbf{Z}|\mathbf{X}, \mathbf{Y}) d\alpha d\theta d\phi,$$

- where

$$p(\alpha, \theta, \phi, \mathbf{Z}|\mathbf{X}, \mathbf{Y}) = \frac{p(\alpha, \theta, \phi, \mathbf{X}, \mathbf{Y}, \mathbf{Z})}{\sum_{\mathbf{Z}} \iiint p(\alpha, \theta, \phi, \mathbf{X}, \mathbf{Y}, \mathbf{Z}) d\alpha d\theta d\phi}.$$

Solution

- Variational algorithm to solve the problem
 - Approximate distribution $p(\alpha, \theta, \phi, \mathbf{Z} | \mathbf{X}, \mathbf{Y})$ by variational distribution $Q = q(\alpha)q(\theta)q(\phi) \prod_i q(Z_i)$
 - $$\begin{aligned}
 \mathbf{Z}^* &= \arg \max_{\mathbf{Z}} p(\mathbf{Z} | \mathbf{X}, \mathbf{Y}) \\
 &= \arg \max_{\mathbf{Z}} \iiint p(\alpha, \theta, \phi, \mathbf{Z} | \mathbf{X}, \mathbf{Y}) d\alpha d\theta d\phi \\
 &\approx \arg \max_{\mathbf{Z}} \iiint q(\alpha, \theta, \phi, \mathbf{Z}) d\alpha d\theta d\phi \\
 &= \arg \max_{\mathbf{Z}} \iiint q(\alpha)q(\theta)q(\phi) \prod_i q(Z_i) d\alpha d\theta d\phi \\
 &= \arg \max_{\mathbf{Z}} \prod_i q(Z_i) \\
 &= \left[\arg \max_{Z_1} q(Z_1), \arg \max_{Z_2} q(Z_2), \dots, \arg \max_{Z_N} q(Z_N) \right]
 \end{aligned}$$
 - The global maximization over \mathbf{Z} reduces to local maximization over each Z_i independently

Variational Approximation

- Approximate $P(Z | X, Y)$ by distribution $Q(Z)$ is to maximize the distribution : $L = \sum_z Q(Z) \log \frac{P(Z, X, Y)}{Q(Z)}$
- $D_{KL}(Q || P) = \sum_z Q(Z) \log \frac{Q(Z)}{P(Z | X, Y)}$
- $D_{KL}(Q || P) = \sum_z Q(Z) \log \frac{Q(Z)}{P(Z, X, Y)} + \log p(X, Y)$
- $\log p(X, Y) = D_{KL}(Q || P) - \sum_z Q(Z) \log \frac{Q(Z)}{P(Z, X, Y)}$
- $\log p(X, Y) = D_{KL}(Q || P) + \sum_z Q(Z) \log \frac{P(Z, X, Y)}{Q(Z)}$

Maximize L

In order to maximize the objective function L, we take the derivatives of L with respect to variational parameters and set these derivatives to zeros.

$$\nabla \tilde{L}(q) = \left(\frac{\partial \tilde{L}}{\partial \tilde{\xi}}, \frac{\partial \tilde{L}}{\partial \tilde{\gamma}}, \frac{\partial \tilde{L}}{\partial \tilde{\mu}}, \frac{\partial \tilde{L}}{\partial \tilde{v}}, \frac{\partial \tilde{L}}{\partial \tilde{\beta}} \right) = 0.$$

Experimental Result

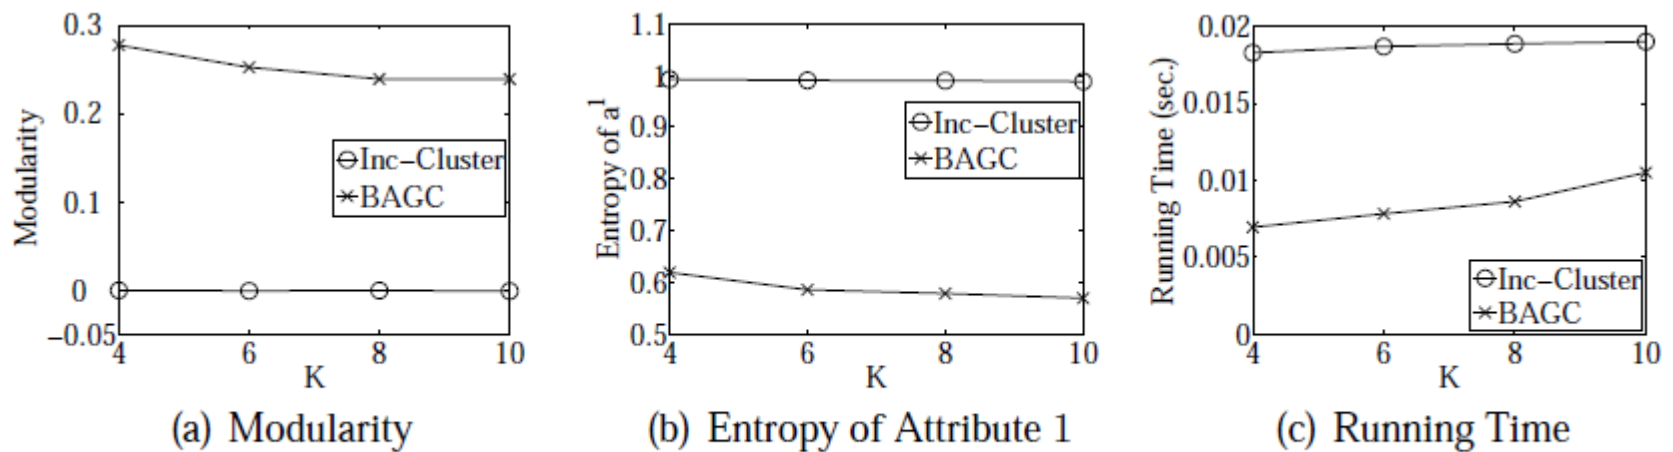


Figure 2: Clustering Performance on Political Blogs

Conclusion

- Introduce the problem of attributed graph clustering and its wide applications
- Introduce two state-of-the-art approaches: distance-based attributed graph clustering and model-based attributed graph clustering
- Demonstrate the performance of both algorithms on the real data set.

Reference

1. [Yang Zhou, Hong Cheng, Jeffrey Xu Yu](#): Graph Clustering Based on Structural/Attribute Similarities. [PVLDB 2\(1\):718-729 \(2009\)](#)
2. A Model-based Approach to Attributed Graph Clustering
3. [Ravindra K. Ahuja, Thomas L. Magnanti, James B. Orlin](#): Network flows - theory, algorithms and applications. [Prentice Hall 1993](#)
4. [Lixin Tao, Yongchang Zhao](#): Multi-way graph partition by stochastic probe. [Computers & OR \(COR\) 20\(3\):321-347 \(1993\)](#)
5. [Matthew J. Rattigan, Marc E. Maier, David Jensen](#): Graph clustering with network structure indices. [ICML 2007:783-790](#)
6. [Charu C. Aggarwal, Haixun Wang](#): A Survey of Clustering Algorithms for Graph Data. [Managing and Mining Graph Data 2010:275-301](#)
7. [Hong Cheng, Yang Zhou, Jeffrey Xu Yu](#): Clustering Large Attributed Graphs: A Balance between Structural and Attribute Similarities. [TKDD 5\(2\):12 \(2011\)](#)
8. [Yang Zhou, Hong Cheng, Jeffrey Xu Yu](#): Clustering Large Attributed Graphs: An Efficient Incremental Approach. [ICDM 2010:689-698](#)

Thank you